Chapter 1

Model

1.1 Coalescent with Local Population Structure

We now present novel local phylodynamic model. In order to help illustrate
the concepts behind this model, consider the following scenario: Suppose there
is a parent population of bacteria of interest evolving through time. At any
point in time, a particular individual within this population gains a significant
evolutionary advantage. This advantage enables the individual and its progeny
to undergo a rapid clonal expansion as it gains the ability to colonise a par-
ticular niche, eventually reaching a constant population size equilibrium. We
will associate the clade corresponding to the clonal expansion with a particular
colouring. The clonal expansion process may repeat throughout time, giving rise
to several clades. The clonal expansions may for example correspond to process
such a bacterial strain gaining antibiotic resistance, or being newly introduced
into a hospital.

In general, we will assume that the clonal expansion process happens relatively
rarely. We will also assume the parent population is at eqillibrium. As un-
der this scenario populations reach a constant population size in the limit this
assumption can be interpreted that the clonal expansion that gave rise to the
parent population has happened a very long time ago.

At present a researcher collects a random set of N samples from the population
of interest, containing at least some of the clades produced by the clonal expan-
sion process described above. Each sample belongs to one of the k clonal expan-
sion clades with multinomial probability θ_i, $\boldsymbol{\theta} = [\theta_1, ..., \theta_i, ..., \theta_{k-1}, \theta_k]$.

The researcher then sequences the genomes of the samples and infers the corre-
sponding phylogenetic using any of the popular available methods. The result-
ing phylogeny can be viewed as a realisation of the following backwards time
process.

Select K clade colours, and an associated clade sampling membership probability vector $\boldsymbol{\theta}$. For each of the N samples s_i forming the leaves of the phylogeny, assign a colouring c_i. The samples of identical colour then coalesce with each other with rate governed by a colour specific population size function $\alpha_i(t)$. As it is assumed each colour is formed by a clonal expansion at time t_i^{Div}, $\alpha_i(t)$ vanishes to zero at the time of the expansion, and as such all coalescence within clade of colour c_i happens almost surely before t_i^{Div}. Upon reaching the time of clonal expansion the MRCA then chnages colour to that of another clade extant at the time. The clade of the colour the MRCA changes to will be referred to as the parent clade. The MRCA is then added to the parent clade as a leaf. The process continues until all only one individual remains.

The coalescent process given leaf colouring and growth function parameters described above is characterised mathematically in the following section.

1.1.1 Model

1.2 Preliminaries

A given genealogy $\mathbf{g} = (V_\mathbf{g}, E_\mathbf{g}, t_\mathbf{g})$ is an incomplete, empirical sample of the underlying process.
It consists of nodes $V_\mathbf{g}$, directed edges $E_\mathbf{g}$, and node labels $t_\mathbf{g}$ corresponding to event times.
The genealogy $\mathbf{g}$ shall be indexed by an index set $S = 1 \leq i \leq N \subset \mathbb{N}$, with $Y \subset S$ corresponding to coalescent events and $I \subset S$ corresponding to sampling events.
For convenience, assume that all edges are in the forwards time direction, i.e.:

$$\forall k, l \in S : (k, l) \in E_\mathbf{g} \Rightarrow t_k < t_l$$

Furthermore, all event times are ordered in descending (backwards) time order, with the first event corresponding the the most recent sample

$$\forall k, l \in S : k < l \Rightarrow t_k > t_l$$

Under the assumption that $\mathbf{g}$ is a genealogy of a given sample, with each edge in $E_\mathbf{g}$ there is an associated unobserved set of individuals descending from one another. At some point along an edge from one lineage to another, the lineage can undergo a colour change, and become the most recent ancestor of a diverging clade. This event corresponds to this lineage somehow gaining advantage over other lineages, be it a bacterium gaining resistance against a drug, or a strain of a virus invading a completely susceptible population.

Definition 1.2.1 (Multiple Lineage Coalescent)**.** Given M colours, M population size functions $\alpha \triangleq \{\alpha_j(t)\}_{1 \leq j \leq M}$. Let $Y(t)$ be a CTMC with the state space:

$$\Sigma = \{\mathbf{s} \in \mathbb{Z}_+ : |\mathbf{s}| \geq 1\} \tag{1.1}$$

and the transition rates

$$\mathbf{s} \to \mathbf{s} - \mathbf{e_j} \qquad \binom{s_j}{2} \alpha_j^{-1}(t) \qquad 1 \le j \le M \qquad (1.2)$$

$$\mathbf{s} \to \mathbf{s} - \mathbf{e_j} + \mathbf{e_k} \qquad \delta_{1,j} \beta s_k \qquad 1 \le j, k \le M \qquad (1.3)$$

Where β is an unknown rate.

The interpretation of this model in backwards (coalescent) time is that each node corresponds to a single specific clade (colour). Nodes of the same clade coalesce at i.i.d rates, according to a clade specific growth functions, until reaching the most recent common ancestor (MRCA) of given clade. The MRCA then changes type (colour) to that of any other clade that is extant at a given time.
Under the hypothesis, along each of the edges from the parent of a clade MRCA to the MRCA lies a point in that characterises the time of divergence of the clade, after which the clade starts to undergo clonal expansion.

1.3 Full Generative Model

Given a population sample of N individuals indexed by F, $X_F = \{x_i\}_{i \in F}$, first determine the number of recent clonal expansions k:

$$k \sim \mathtt{poi}(\phi) \qquad (1.4)$$

Then simulate the $k + 1$ clonal expansion probability vector $\boldsymbol{\theta}$:

$$\boldsymbol{\theta} \sim \mathtt{dirichlet}(\psi), \quad \psi \in \mathbb{R}_+^{k+1} \qquad (1.5)$$

Given the number of expansions k and expansion probabilities $\boldsymbol{\theta}$, partition F into $k + 1$ mutually disjoint subsets $\mathbf{f} = \{f_1, ... f_k, f_{k+1}\}$ with $\bigcup\limits_{i=1}^{k+1} f_i = F$ with the probability $P[j \in f_i] = \theta_i \quad \forall i, \forall j \in F$. Each partition corresponds to an individual colour characterising a clonal expansion. For convenience, we will set f_{k+1} to be the index set of the tips corresponding to the ancestral subpopulation.

Individual subpopulations then follow the coalescent process described above in 1.2.1, governed by population size functions $\alpha_j^{-1}(t)$. To each of α_j^{-1} corresponds a set of parameters:

$$
\begin{aligned}
T_{div_j} \in \mathbb{R}_+ &\quad \text{The time of divergence} \\
r_j \in \mathbb{R}_+ &\quad \text{The growth rate of the subpopulation} \qquad (1.6) \\
N_j \in \mathbb{R}_+ &\quad \text{The carrying capacity}
\end{aligned}
$$

Upon reaching time of divergence, a subpopulation then changes colour to that of its parent subpopulation ρ_i. ρ_i are selected uniformly at random from all extant populations at time T_{div_i}.

In the case of the ancestral subpopulation f_{k+1}, we assume that the divergence event happened a very long time ago. As by definition $\alpha(t) \to N$ as $t - T_{div} \to -\infty$, we approximate $\alpha_{k+1}(t) \approx N_{k+1}$ and as such $\alpha_{k+1}^{-1}(t) \approx 1/N_{k+1}$. The parameters T_{div_j}, r_j, N_j are simulated from appropriate (prior) distributions. Finally, without loss of generality, assume that divergence events are indexed in descending order by the time of divergence.

$$i > j \Leftrightarrow T_{div_i} < T_{div_j} \tag{1.7}$$

Finally, simulate the sampling times $\mathbf{t} = \{t_i\}_{i \in F}$ such that for all expansions j, $t_i < T_{div_j}$, for all $i \in f_j$. In other words, all sampling for a given clade must happen after the clade diverges.

1.3.1   Posterior

The posterior for the full process is:

$$\begin{aligned}
P(k, \mathbf{f}, \boldsymbol{\theta}, \boldsymbol{T_{div}}, \mathbf{r}, \mathbf{N} \mid \mathbf{g}) \propto\; & \mathcal{L}(\mathbf{g} \mid \mathbf{f}, \boldsymbol{\rho}, \boldsymbol{T_{div}}, \mathbf{r}, \mathbf{N}) \\
& \times \mathcal{L}(\mathbf{f} \mid \boldsymbol{\theta}, k) \\
& \times \pi(\boldsymbol{\rho} \mid \mathbf{T_{div}}) \\
& \times \pi(\mathbf{T_{div}}, \mathbf{r}, \mathbf{N} \mid k) \\
& \times \pi(\boldsymbol{\theta} \mid k)\pi(k)
\end{aligned} \tag{1.8}$$

The first likelihood term

$$\mathcal{L}(\mathbf{g} \mid \mathbf{f}, \boldsymbol{\rho}, \boldsymbol{T_{div}}, \mathbf{r}, \mathbf{N}) \tag{1.9}$$

corresponds to the likelihood of the coalescent with local population strucuture described in 1.2.1. It can be expressed as a product of likelihoods of clonal expansion subtrees with diverging clonal expansions added as leaves under varying population size coalescent process given the corresponding population size function α_i.

$$\mathcal{L}(\mathbf{g} \mid \mathbf{f}, \boldsymbol{\rho}, \boldsymbol{T_{div}}, \mathbf{r}, \mathbf{N}) = \prod_{i=1}^{k+1} \mathcal{L}(\mathbf{g}_i \mid T_{div_i}, r_i, N_i) \tag{1.10}$$

Where $\mathbf{g_i}$ is the clonal expansion subtree corresponding to expansion i, contanining all the leaves in f_i, the entire clade specified by theses, as well as any divergence events j added as leaves if $\rho_j = i$. The second term in the likelihood

$$\mathcal{L}(\mathbf{f} \mid \boldsymbol{\theta}, k) \tag{1.11}$$

Corresponds to the likelihood of samples corresponding to given clonal expansion clades, given the clonal expansion clade sampling probabilities. It is simply a multinomial probability

$$\mathcal{L}(\mathbf{f} \mid \boldsymbol{\theta}, k) = \prod_{i=1}^{k} \theta_i^{|f_i|} \tag{1.12}$$

As the model dimensionality varies with k, the posterior has to be integrated via RJ-MCMC. Furthermore as $\boldsymbol{\theta}$ has to be a multinomial probability vector, it must be rescaled with every transdimensional move so that it sums to one, leading to non-unitary determinants. For inference we use the following priors: The prior on the parental clades of individual clonal expansions

$$\pi(\boldsymbol{\rho} \mid \mathbf{T_{div}}) \tag{1.13}$$

is assumed to be uniform probability that the clonal expansion diverged from any other subpopulation extant at the time, i.e.:

$$\pi(\boldsymbol{\rho} \mid \mathbf{T_{div}}) \sim \prod_{i=1}^{k} \mathbb{1}_{\{x:x>i\}}(\rho_i)\frac{1}{k+1-i} \tag{1.14}$$

The prior on functions of effective population size can be expressed as

$$\pi(\mathbf{T_{div}}, \mathbf{r}, \mathbf{N} \mid k) \sim \pi(N_{k+1}) \prod_{i=1}^{k} \pi(N_i)\pi(T_{div_i})\pi(r_i) \tag{1.15}$$

For $\pi(N_i) \quad \forall i \in \{0, ..., k-1, k\}$, log-normal distributions parametrised by mean μ_N and standard deviation σ_N provided by the decision maker are used.
$\pi(r_i) \quad \forall i \in \{1, ..., k-1, k\}$, log-normal distributions parametrised by mean μ_r and standard deviation σ_r provided by the decision maker are used.
$\pi(T_{div_i})$ are assumed to be gamma distributed, parametrised by mean and variance supplied by the decision maker.
For prior on membership probabilities

$$\pi(\boldsymbol{\theta} \mid k) \sim \texttt{dirichlet}(\psi) \tag{1.16}$$

k-dimensional dirichlet distribution is used, with ψ supplied by the decision maker.

For prior on the number of clonal expansions

$$\pi(k) \sim \texttt{poi}(1) \tag{1.17}$$

we use poisson distribution with rate 1

Chapter 2

References